

CDK/Ceph installation via MaaS and Juju

Final components layout:

	Bare Metal	Container #1	Container #2	Container #3	Container #4	Container #5	VM #1
Node #1	Ceph OSD	Ceph mon	k8s master/ flannel	etcd	flannel	easyrsa	k8s worker/ flannel
Node #2	Ceph OSD	Ceph mon	k8s master/ flannel	etcd	flannel	RADOSGW	k8s worker/ flannel
Node #3	Ceph OSD	Ceph mon	k8s master/ flannel	etcd	flannel	kubeapi load balancer	k8s worker/ flannel
	Bare Metal	VM #1					
Infra #1	MAAS Rack + Region	Juju Controller					

Pre-requisites:

- Have a VM created on the infra node with a pre-defined name (let it be “juju-controller”)

For that:

Adjust netplan config for having juju-controller VM to be connected to the bridge:

```
network:
  ethernets:
    enp129s0f0:
      addresses: []
      dhcp4: true
    enp129s0f1:
      addresses: []
      dhcp4: true
    enp1s0f0:
      dhcp4: no
      addresses: [172.17.1.6/16]
      gateway4: 172.17.0.1
      nameservers:
        addresses: [ 8.8.8.8,8.8.4.4 ]
    enp1s0f1:
      match:
        macaddress: 0c:c4:7a:f7:ef:71
    enp1s0f2:
      addresses: []
```

```

        dhcp4: true
        optional: true
    enpls0f3:
        addresses: []
        dhcp4: true
        optional: true
    bridges:
        br0:
            dhcp4: no
            addresses: [192.168.1.1/24]
            interfaces:
                - enpls0f1
    version: 2

```

After that run

```

sudo netplan try
#in case network is functioning properly - accept changes

```

Set up masquerade routing for the nodes:

```

smadmin@sm-maas-kb:~$ cat iptables
#!/bin/bash
iptables -t nat -A POSTROUTING -o enpls0f0 -j MASQUERADE
iptables -A FORWARD -i enpls0f0 -o br0 -m state --state RELATED,ESTABLISHED -j
ACCEPT
iptables -A FORWARD -i enpls0f0 -o br0 -j ACCEPT
smadmin@sm-maas-kb:~$ chmod +x iptables
smadmin@sm-maas-kb:~$ sudo ./iptables

```

Create a VM using virt-install:

```

virt-install --name=juju-controller --disk size=50,sparse=no,pool=images --virt-type
kvm --graphics spice --vcpus=2 --ram=4096 --pxe --network bridge=br0
--os-type=linux --os-variant=ubuntu16.04

```

- MaaS should have 2 network spaces defined, one for general purposes “**default**” and another one for Ceph replication “**ceph-cluster**”. This can be done via MaaS UI, section “Subnets”:

The screenshot shows the MAAS Subnets page. At the top, there's a navigation bar with 'MAAS' and tabs for 'Nodes', 'Pods', 'Images', 'DNS', 'Zones', 'Subnets', and 'Settings'. The 'Subnets' tab is active. Below the navigation bar, there's a 'Subnets' header and a 'Group by' dropdown set to 'Fabrics'. A table lists subnets with columns: Fabric, VLAN, DHCP, Subnet, Available IPs, and Space. An 'Add' dropdown menu is open, showing options: Fabric, VLAN, Space, and Subnet.

Fabric	VLAN	DHCP	Subnet	Available IPs	Space
fabric-0	untagged	Enabled	192.168.122.0/24	63%	default
	1 (ceph)	Disabled	172.16.0.0/24	99%	ceph-cluster

MAAS name: maas1 MAAS
MAAS version: 2.3.0 (6436-gd354690-ubuntu1-16.04.1)
[View release notes](#) · [View documentation](#)
© 2018 Canonical Ltd. Ubuntu and Canonical are registered trademarks of Canonical Ltd.

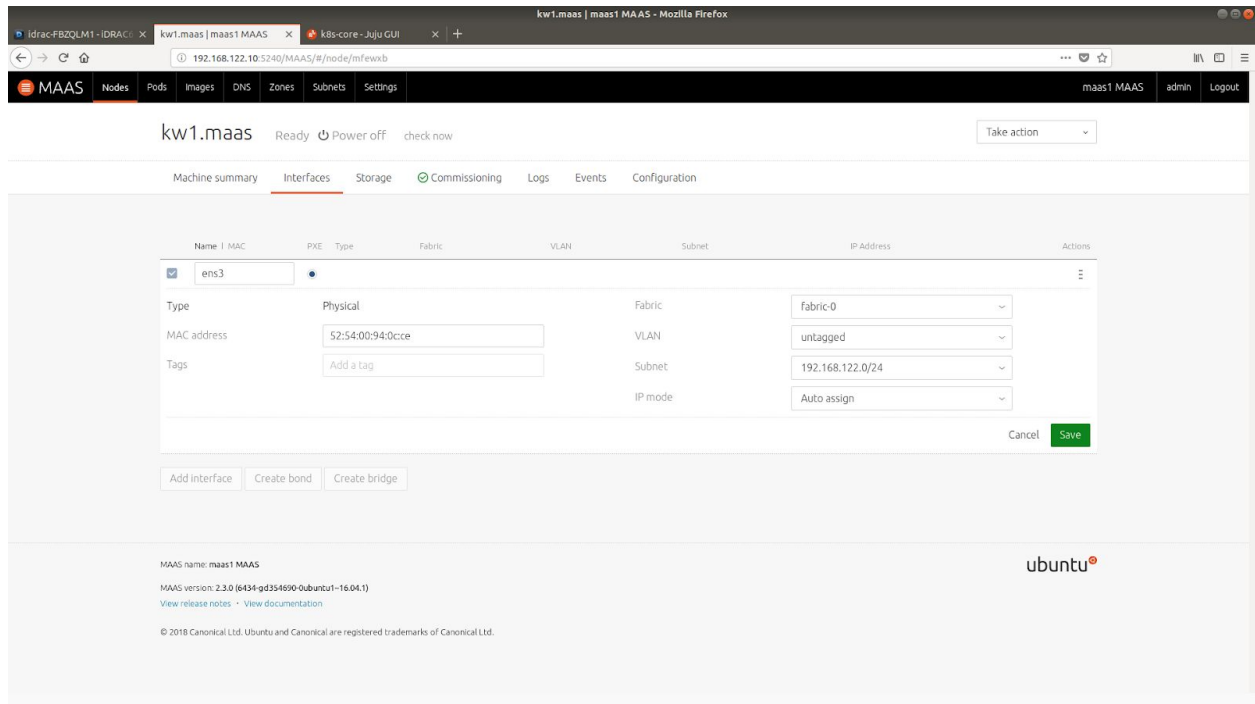
After that each of 2 subnets should be edited to belong to its own space.

The screenshot shows the MAAS Subnet summary page for the subnet 192.168.122.0/24. The page has a 'Take action' dropdown at the top right. Below the header, there's a 'Subnet summary' section with an 'Edit' button. The summary table shows details for the subnet, including Name, CIDR, Gateway IP, DNS, Description, Managed allocation, Active discovery, Fabric, VLAN, and Space. Below the summary, there's a 'Static Routes' section with a table for Gateway IP, Destination, Metric, and Actions. A message states 'No static routes have been added to this subnet.' and there's an 'Add static route' button.

Name	CIDR	Gateway IP	DNS	Description	Managed allocation	Active discovery	Fabric	VLAN	Space
192.168.122.0/24	192.168.122.0/24	192.168.122.1			Enabled	Disabled	fabric-0	untagged	default

Gateway IP	Destination	Metric	Actions
No static routes have been added to this subnet.			

- All cloud nodes should have their interfaces configured - assigned to the proper VLAN and subnet, IP address in “Auto-assigned” mode.



All cloud nodes should be in “Ready” state.

!NOTE Do not use LVM for the OS disks, otherwise nodes boot will never happen.

Further actions should happen on the separate machine/container you will treat as your juju client (Ubuntu 18.04)

```
sudo snap install juju --classic
```

Get oauth1 credential from MaaS (API key, in UI under “admin”, MaaS keys)

```
ubuntu@ubuntu:~$ juju add-cloud
Cloud Types
  maas
  manual
  openstack
  oracle
  vsphere

Select cloud type: maas

Enter a name for your maas cloud: cloud1

Enter the API endpoint url: http://192.168.122.10:5240/MAAS

Cloud "cloud1" successfully added
You may bootstrap with 'juju bootstrap cloud1'

ubuntu@ubuntu:~$ juju add-credential cloud1
```

Enter credential name: cloud1-maas-creds

Using auth-type "oauth1".

Enter maas-oauth:

Credential "cloud1-maas-creds" added locally for cloud "cloud1".

ubuntu@ubuntu:~\$ juju bootstrap cloud1 controller1 --to=juju-controller

Creating Juju controller "controller1" on cloud1

Looking for packaged Juju agent version 2.4.2 for amd64

Launching controller instance(s) on cloud1...

- wh4wse (arch=amd64 mem=4G cores=2)

Installing Juju agent on bootstrap instance

Fetching Juju GUI 2.13.2

Waiting for address

Attempting to connect to 192.168.122.12:22

Connected to 192.168.122.12

Running machine configuration script...

Bootstrap agent now started

Contacting Juju controller at 192.168.122.12 to verify accessibility...

Bootstrap complete, "controller1" controller now available

Controller machines are in the "controller" model

Initial model "default" added

ubuntu@ubuntu:~\$ juju add-model cdk-ceph

Uploading credential 'cloud1/admin/cloud1-maas-creds' to controller

Added 'cdk-ceph' model with credential 'cloud1-maas-creds' for user 'admin'

ubuntu@ubuntu:~\$ juju deploy cdk-ceph-small.yaml

...

ubuntu@ubuntu:~\$ juju status

Model	Controller	Cloud/Region	Version	SLA	Timestamp
k8s-core	controller1	cloud1	2.4.2	unsupported	21:57:04Z

App	Version	Status	Scale	Charm	Store	Rev	OS
Notes							
easysrsa		waiting	1/2	easysrsa	jujucharms	68	
ubuntu							
etcd	3.2.9	blocked	1	etcd	jujucharms	126	
ubuntu							
flannel	0.10.0	blocked	2	flannel	jujucharms	81	
ubuntu							
kubernetes-master	1.11.2	active	1	kubernetes-master	jujucharms	144	
ubuntu exposed							
kubernetes-worker	1.11.2	waiting	1	kubernetes-worker	jujucharms	163	
ubuntu exposed							

Unit	Workload	Agent	Machine	Public address	Ports
Message					
easysrsa/0*	maintenance	executing	0/lxd/0	192.168.122.15	
(install) installing charm software					
easysrsa/1	waiting	allocating	1/lxd/0		
waiting for machine					
etcd/0*	blocked	idle	0	192.168.122.13	
Missing relation to certificate authority.					
kubernetes-master/0*	active	idle	0	192.168.122.13	
Kubernetes master running.					

```

flannel/0*          blocked      idle          192.168.122.13
Waiting for etcd relation.
kubernetes-worker/0* waiting      idle          1          192.168.122.14
Waiting for cluster DNS.
flannel/1           maintenance  idle          192.168.122.14
Unpacking flannel resource.

Machine  State    DNS           Inst id          Series  AZ          Message
0        started  192.168.122.13 etke74           bionic  default     Deployed
0/lxd/0  started  192.168.122.15 juju-3f7cb5-0-lxd-0 bionic  default     Container
started
1        started  192.168.122.14 yf7y6p           bionic  default     Deployed
1/lxd/0  pending                juju-3f7cb5-1-lxd-0 bionic  default     Container
started

```

Once the deployment is done there should be k8s cluster created on top of 4 machines.

```

ubuntu@ubuntu:~$ mkdir -p ~/.kube
ubuntu@ubuntu:~$ juju scp kubernetes-master/0:config ~/.kube/config
ubuntu@ubuntu:~$ sudo snap install kubectl --classic
[sudo] password for ubuntu:
kubectl 1.11.2 from Canonical✓ installed
ubuntu@ubuntu:~$ kubectl cluster-info
Kubernetes master is running at https://192.168.122.13:6443
Heapster is running at
https://192.168.122.13:6443/api/v1/namespaces/kube-system/services/heapster/proxy
KubeDNS is running at
https://192.168.122.13:6443/api/v1/namespaces/kube-system/services/kube-dns:dns/proxy
Kubernetes-dashboard is running at
https://192.168.122.13:6443/api/v1/namespaces/kube-system/services/https:kubernetes-
dashboard:/proxy
Metrics-server is running at
https://192.168.122.13:6443/api/v1/namespaces/kube-system/services/https:metrics-ser
ver:/proxy
Grafana is running at
https://192.168.122.13:6443/api/v1/namespaces/kube-system/services/monitoring-grafan
a/proxy
InfluxDB is running at
https://192.168.122.13:6443/api/v1/namespaces/kube-system/services/monitoring-influx
db:http/proxy

To further debug and diagnose cluster problems, use 'kubectl cluster-info dump'.

```

If it is necessary to extend the cluster (add physical worker node):

```

juju add-unit ceph-osd
juju add-unit kubernetes-worker --to=kvm:<new_machine_id>

```

Accessing Kubernetes cluster

1. Install kubectl on your machine (for Ubuntu - "sudo snap install kubectl --classic")

Create directories:

```
mkdir -p ~/.kube
```

2. Copy kube config from MaaS server to your machine:

```
scp -r smadmin@64.169.30.89:~/.kube/* /home/agrebennikov/.kube/
```

3. Adjust the address of the k8s API server:

```
sed -i 's/192.168.2.100:443/127.0.0.1:8443/' ~/.kube/config
```

4. Establish a tunnel to the MaaS server with port forwarding:

```
ssh -L 8443:192.168.2.100:443 smadmin@64.169.30.89
```

Make sure nobody else is occupying the same port 8443 at the moment, and if so - use another dynamic port.

5. Start kube proxy on the local machine:

```
kubectl proxy
```

6. In the browser access the cluster via URL:

```
http://localhost:8001/api/v1/namespaces/kube-system/services/https:kubernetes-dashboar:~board:/proxy/
```

Appendix 1

CDK+Ceph bundle

Marked in red is unique per cluster and needs to be adjusted.

```
series: bionic
machines:
  '0':
    series: bionic
  '1':
    series: bionic
  '2':
    series: bionic
applications:
  keepalived:
    charm: cs:~containers/keepalived-4
    annotations:
      gui-x: '450'
      gui-y: '750'
    options:
      virtual_ip: 192.168.1.100/24
  flannel:
    charm: cs:~containers/flannel
    annotations:
      gui-x: '450'
      gui-y: '750'
  flannel-worker:
    charm: cs:~containers/flannel
    annotations:
      gui-x: '450'
      gui-y: '750'
```

```
options:
  iface: enp0s2
easyrsa:
  charm: cs:~containers/easyrsa
  num_units: 1
  annotations:
    gui-x: '450'
    gui-y: '550'
  to:
    - lxd:ceph-osd
  bindings:
    "": default
ceph-osd:
  charm: cs:ceph-osd
  num_units: 3
  options:
    osd-devices: /dev/nvme0n1 /dev/nvme1n1
  bindings:
    "": default
    cluster: ceph-storage
  annotations:
    gui-x: '300'
    gui-y: '300'
  to:
    - 0
    - 1
    - 2
ceph-radosgw:
  annotations:
    gui-x: '1000'
    gui-y: '250'
  charm: cs:ceph-radosgw
  num_units: 1
  bindings:
    "": default
  to:
    - lxd:ceph-osd
kubernetes-worker:
  charm: cs:~containers/kubernetes-worker
  constraints: cores=24 mem=393216 root-disk=2560G
  num_units: 3
  expose: true
  annotations:
    gui-x: '100'
    gui-y: '850'
  bindings:
    "": default
  to:
    - 'kvm:0'
    - 'kvm:1'
    - 'kvm:2'
kubernetes-master:
  charm: cs:~containers/kubernetes-master
  num_units: 3
  annotations:
    gui-x: '800'
    gui-y: '850'
  bindings:
    "": default
  to:
```



```
- lxd:0
- lxd:1
- lxd:2
ceph-mon:
  charm: 'cs:ceph-mon'
  num_units: 3
  annotations:
    gui-x: '600'
    gui-y: '300'
  bindings:
    "": default
  to:
    - lxd:0
    - lxd:1
    - lxd:2
etcd:
  charm: cs:~containers/etcd
  num_units: 3
  annotations:
    gui-x: '800'
    gui-y: '550'
  bindings:
    "": default
  to:
    - lxd:ceph-osd
kubeapi-load-balancer:
  charm: cs:~containers/kubeapi-load-balancer
  num_units: 1
  expose: true
  annotations:
    gui-x: '450'
    gui-y: '250'
  bindings:
    "": default
  options:
    extra_sans: 192.168.1.100
  to:
    - lxd:ceph-osd
relations:
- - 'ceph-mon:osd'
- - 'ceph-osd:mon'
- - 'kubernetes-master:kube-api-endpoint'
- - 'kubeapi-load-balancer:apiserver'
- - 'kubernetes-master:loadbalancer'
- - 'kubeapi-load-balancer:loadbalancer'
- - 'kubernetes-master:kube-control'
- - 'kubernetes-worker:kube-control'
- - 'kubernetes-master:certificates'
- - 'easyrsa:client'
- - 'etcd:certificates'
- - 'easyrsa:client'
- - 'kubernetes-master:etcd'
- - 'etcd:db'
- - 'kubernetes-worker:certificates'
- - 'easyrsa:client'
- - 'kubernetes-worker:kube-api-endpoint'
- - 'keepalived:website'
- - 'kubernetes-master:loadbalancer'
- - 'keepalived:loadbalancer'
- - 'kubeapi-load-balancer:website'
```

```
- 'keepalived:lb-sink'
- - 'kubeapi-load-balancer:juju-info'
- - 'keepalived:juju-info'
- - 'kubeapi-load-balancer:certificates'
- - 'easyrsa:client'
- - 'flannel:etcd'
- - 'etcd:db'
- - 'flannel-worker:etcd'
- - 'etcd:db'
- - 'flannel:cni'
- - 'kubernetes-master:cni'
- - 'flannel-worker:cni'
- - 'kubernetes-worker:cni'
- - 'kubernetes-master:ceph-storage'
- - 'ceph-mon:admin'
- - 'ceph-mon:radosgw'
- - 'ceph-radosgw:mon'
```

Marked in red is unique per cluster and needs to be adjusted.

Issues

to activate Pod Security Policy:

```
ADMISSION_CONTROLLERS="$(juju ssh kubernetes-master/0 sudo snap get kube-apiserver
admission-control)"
```

```
juju config kubernetes-master
```

```
api-extra-args=admission-control=PodSecurityPolicy,$ADMISSION_CONTROLLERS
```

and then you need to create policies on top of that, as described in

<https://kubernetes.io/docs/concepts/policy/pod-security-policy/>, to allow pods to be created